

# Other Constraints and Assertions

# “check” constraints

- We’ve seen a check clause on a user-defined domain:

```
create domain Grade as smallint  
    default null  
    check (value >= 0 and value <= 100);
```

- You can also define a check constraint
  - on an attribute
  - on the tuples of a relation
  - across relations

# Attribute-based “check” constraints

- Defined with a single attribute and constrain its value (in every tuple).
- Can only refer to that attribute.
- Can include a subquery.

- Example:

```
create table Student (  
    sID integer,  
    program varchar(5) check  
        (program in (select post from P)),  
    firstName varchar(15) not null, ...);
```

- Condition can be anything that could go in a **WHERE** clause.

# When they are checked

- Only when a tuple is inserted into that relation, or its value for that attribute is updated.
- If a change somewhere else violates the constraint, the DBMS will not notice. E.g.,
  - If a student's program changes to something not in table P, we get an error.
  - But if table P drops a program that some student has, there is no error.

# “not null” constraints

- You can declare that an attribute of a table is NOT NULL.

```
create table Course(  
    cNum integer,  
    name varchar(40) not null,  
    dept Department,  
    wr boolean,  
    primary key (cNum, dept));
```

- In practise, many attributes should be `not null`.
- This is a very specific kind of attribute-based constraint.

# Tuple-based “check” constraints

- Defined as a separate element of the table schema, so can refer to *any* attributes of the table.
- Again, condition can be anything that could go in a **WHERE** clause, and can include a subquery.

- **Example:**

```
create table Student (  
    sID integer,  
    age integer, year integer,  
    college varchar(4),  
    check (year = age - 18),  
    check college in  
        (select name from Colleges));
```

# When they are checked

- Only when a tuple is inserted into that relation, or updated.
- Again, if a change somewhere else violates the constraint, the DBMS will not notice.

# How nulls affect “check” constraints

- A check constraint only fails if it evaluates to false.
- It is not picky like a WHERE condition.
- E.g.: `check (age > 0)`

| age  | Value of condition | CHECK outcome | WHERE outcome |
|------|--------------------|---------------|---------------|
| 19   | TRUE               | pass          | pass          |
| -5   | FALSE              | fail          | fail          |
| NULL | unknown            | pass          | fail          |



# Example

- Suppose you created this table:

```
create table Frequencies(  
    word varchar(10),  
    num integer,  
    check (num > 5));
```

- It would allow you to insert ('hello', null)  
since null passes the constraint check (num > 5)
- If you need to prevent that, use a “not null”  
constraint.

```
create table Frequencies(  
    word varchar(10),  
    num integer not null,  
    check (num > 5));
```

# Naming your constraints

- If you name your constraint, you will get more helpful error messages.
- This can be done with any of the types of constraint we've seen.

- Add

`constraint «name»`

before the

`check ( «condition» )`

# Examples

```
create domain Grade as smallint
    default null
    constraint gradeInRange
        check (value >= 0 and value <= 100));
```

```
create domain Campus as varchar(4)
    not null
    constraint validCampus
        check (value in ('StG', 'UTM', 'UTSC'));
```

```
create table Offering(...
    constraint validCourseReference
        foreign key (cNum, dept) references Course);
```

- Order of constraints doesn't matter, and doesn't dictate the order in which they're checked.

# Assertions

- Check constraints can't express complex constraints across tables, e.g.,
  - Every loan has at least one customer, who has an account with at least \$1,000.
  - For each branch, the sum of all loan amounts < the sum of all account balances.
- Assertions are schema elements at the top level, so can express cross-table constraints:  
`create assertion (<name>) check (<predicate>);`

# Powerful but costly

- SQL has a fairly powerful syntax for expressing the predicates, including quantification.
- Assertions are costly because
  - They have to be checked upon every database update (although a DBMS may be able to limit this).
  - Each check can be expensive.
- Testing and maintenance are also difficult.
- So assertions must be used with great care.

# Triggers

- Assertions are powerful, but costly.
- Check constraints are less costly, but less powerful.
- Triggers are a compromise between these extremes:
  - They are cross-table constraints, as powerful as assertions.
  - But you control the cost by having control over when they are applied.

# The basic idea

- You specify three things.
  - Event: Some type of database action, e.g.,  
after delete on Courses or  
before update of grade on Took
  - Condition: A boolean-valued expression, e.g.,  
when grade > 95
  - Action: Any SQL statements, e.g.,  
insert into Winners values (sID)



# Reaction Policies

# Example

- Suppose  $R = \text{Took}$  and  $S = \text{Student}$ .
- What sorts of action must simply be rejected?
- But a deletion or update with an sID that occurs in  $\text{Took}$  could be allowed ...

# Possible policies

- `cascade`: propagate the change to the referring table
- `set null`: set the referring attribute(s) to null
- There are other options we won't cover.  
Many DBMSs don't support all of them.
- If you say nothing, the default is to forbid the change in the referred-to table.

# Reaction policy example

- In the University schema, what should happen in these situations:
  - csc343 changes number to be 543
  - student 99132 is deleted
  - student 99132's grade in csc148 is raised to 85.
  - csc148 is deleted

# Note the asymmetry

- Suppose table R refers to table S.
- You can define “fixes” that propagate changes backwards from S to R.
- (You define them in table R because it is the table that will be affected.)
- You cannot define fixes that propagate forward from R to S.

# Syntax for specifying a reaction policy

- Add your reaction policy where you specify the foreign key constraint.

- Example:

```
create table Took (  
    ...  
    foreign key (sID) references Student  
        on delete cascade  
    ...  
);
```

# What you can react to

- Your reaction policy can specify what to do either
  - `on delete`, i.e., when a deletion creates a dangling reference,
  - `on update`, i.e., when an update creates a dangling reference,
  - or both. Just put them one after the other.

Example:

`on delete restrict on update cascade`

# What your reaction can be

- Your policy can specify one of these reactions (there are others):
  - `restrict`: Don't allow the deletion/update.
  - `cascade`: Make the same deletion/update in the referring tuple.
  - `set null`: Set the corresponding value in the referring tuple to null.



# Semantics of Deletion

- What if deleting a tuple violates a foreign key constraint?

- Example:

```
DELETE FROM Course  
WHERE dept = 'CSC';
```

# Semantics of Deletion

- What if deleting one tuple violates a foreign key constraint, but deleting others does not?

# DDL Wrap-up

# Updating the schema itself

- Alter: alter a domain or table

```
alter table Course
```

```
    add column numSections integer;
```

```
alter table Course
```

```
    drop column breadth;
```

- Drop: remove a domain, table, or whole schema

```
drop table course;
```

- How is that different from this?

```
delete from course;
```

- If you drop a table that is referenced by another table, you must specify “cascade”

- This removes all referring rows.

# There's more to DDL

- For example, you can also define:
  - indices: for making search faster (we'll discuss these later).
  - privileges: who can do what with what parts of the database
- See csc443.